

# Multi-Version Concurrency Control



**Lecture #20**



**Database Systems**

15-445/15-645

Fall 2017



**Andy Pavlo**

Computer Science Dept.  
Carnegie Mellon Univ.

# ADMINISTRIVIA

**Homework #4:** Monday November 13<sup>th</sup> @ 11:59pm

**Project #3:** Wednesday November 15<sup>th</sup> @ 11:59am

# MULTI-VERSION CONCURRENCY CONTROL

The DBMS maintains multiple physical versions of a single logical object in the database:

- When a txn writes to an object, the DBMS creates a new version of that object.
- When a txn reads an object, it reads the newest version that existed when the txn started.



# MULTI-VERSION CONCURRENCY CONTROL

Protocol was first proposed in 1978 MIT PhD dissertation.

First implementation was InterBase (now Firebird).

→ Implemented by Jim Starkey, co-founder of NuoDB.

MVCC is now used in almost every new DBMS implemented in last 10 years.



# MULTI-VERSION CONCURRENCY CONTROL


Writers don't block readers.  
Readers don't block writers.

Read-only txns can read a consistent  
**snapshot** without acquiring locks.  
→ Use timestamps to determine visibility.

Easily support **time-travel queries**.




# MVCC – EXAMPLE #1



## Database

Version	Value	Begin	End
$A_0$	123	0	-

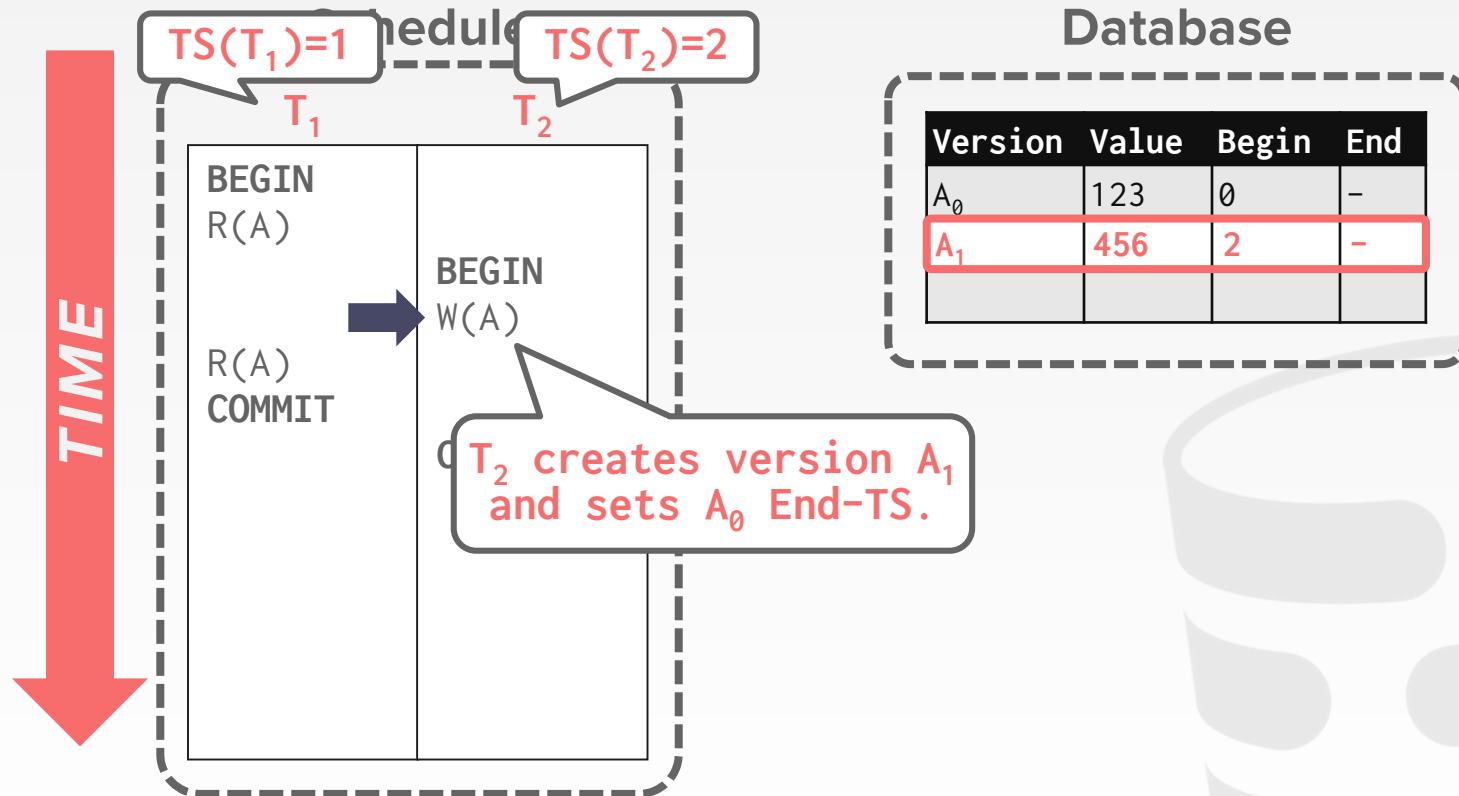
# MVCC – EXAMPLE #1



## Database

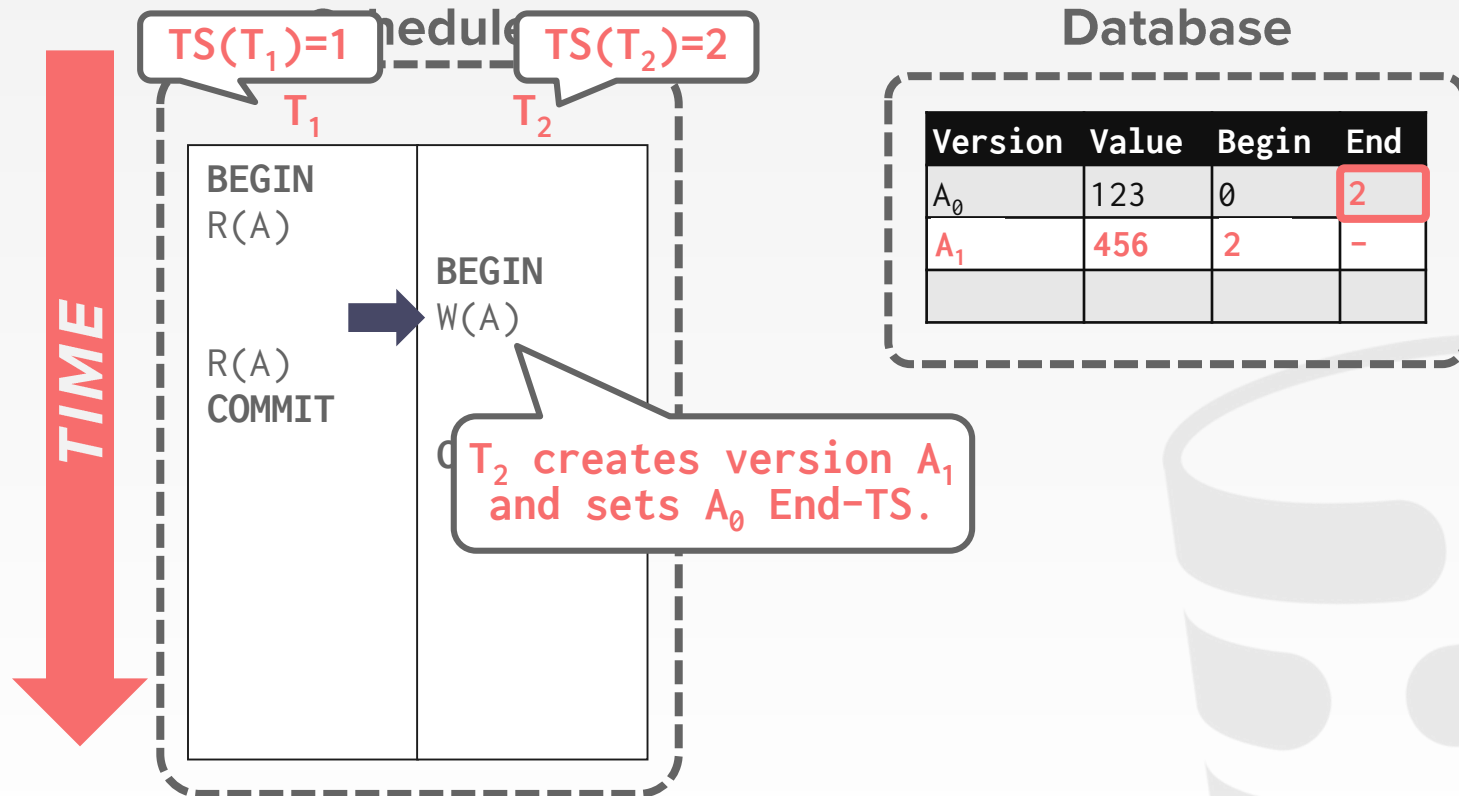
Version	Value	Begin	End
$A_0$	123	0	-

# MVCC – EXAMPLE #1






# MVCC – EXAMPLE #1




# MVCC – EXAMPLE #1



## Database

Version	Value	Begin	End
$A_0$	123	0	2
$A_1$	456	2	-


# MVCC – EXAMPLE #2



## Database

Version	Value	Begin	End
$A_0$	123	0	


# MVCC – EXAMPLE #2



## Database

Version	Value	Begin	End
$A_0$	123	0	
$A_1$	456	1	-

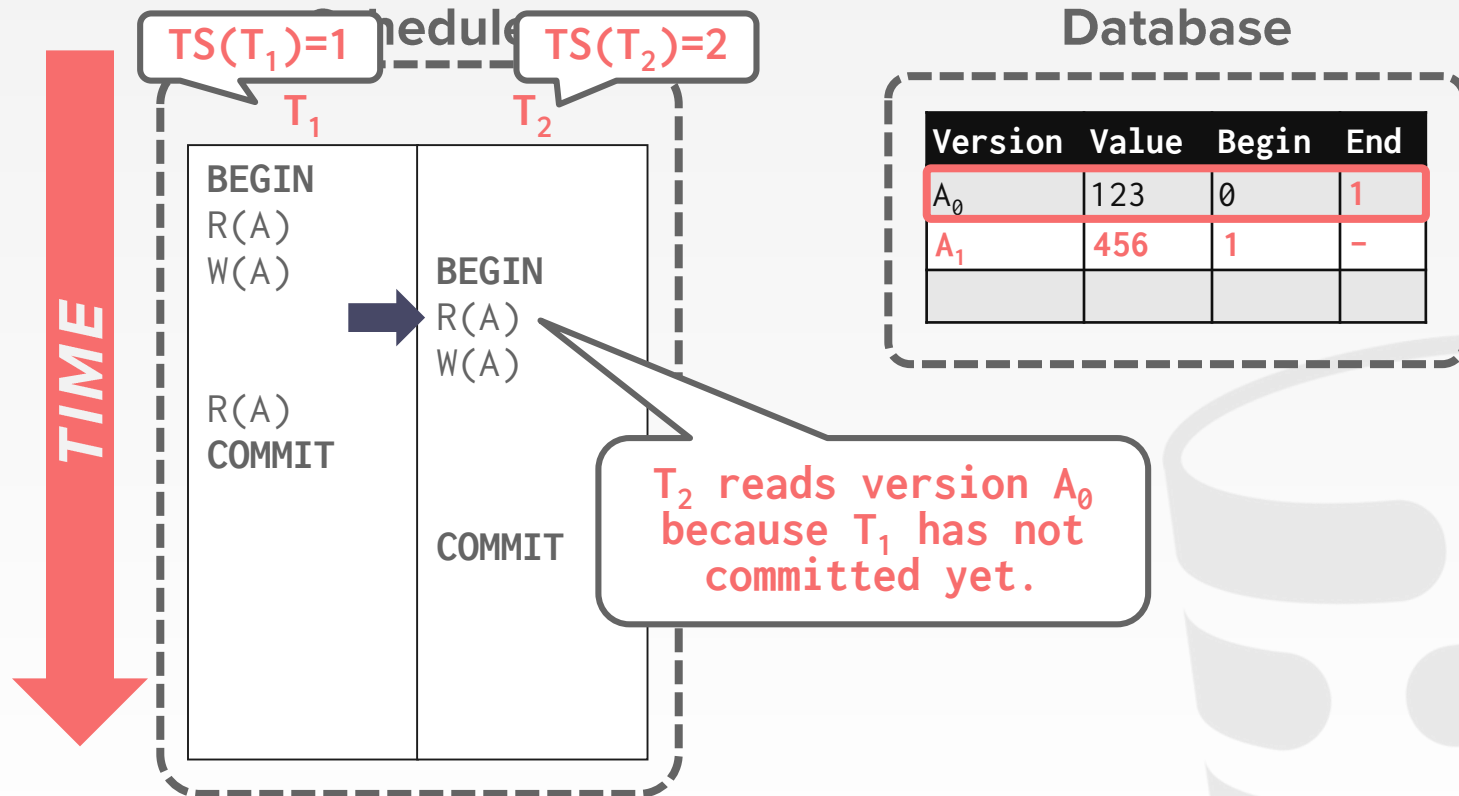
# MVCC – EXAMPLE #2



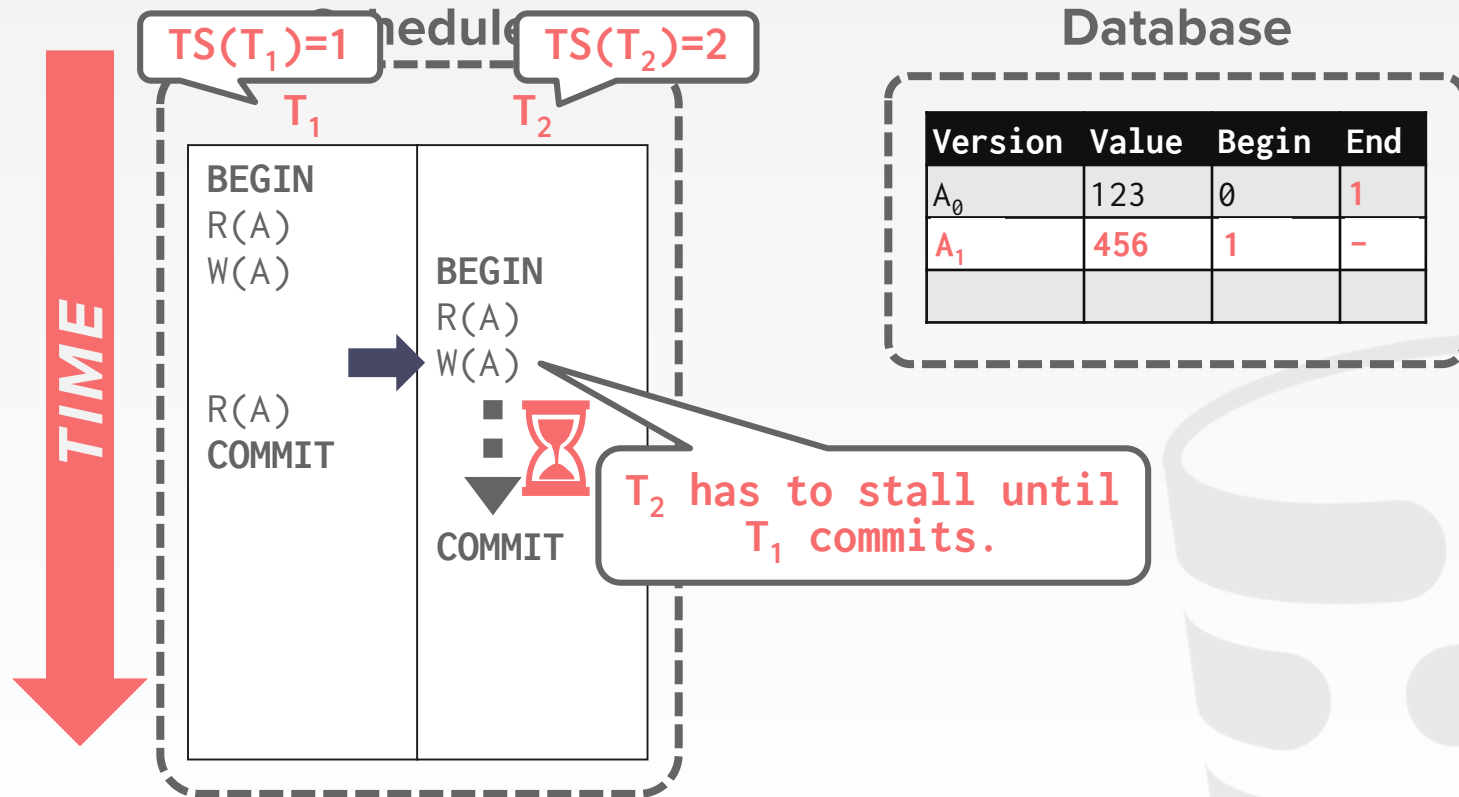
## Database

Version	Value	Begin	End
$A_0$	123	0	1
$A_1$	456	1	-

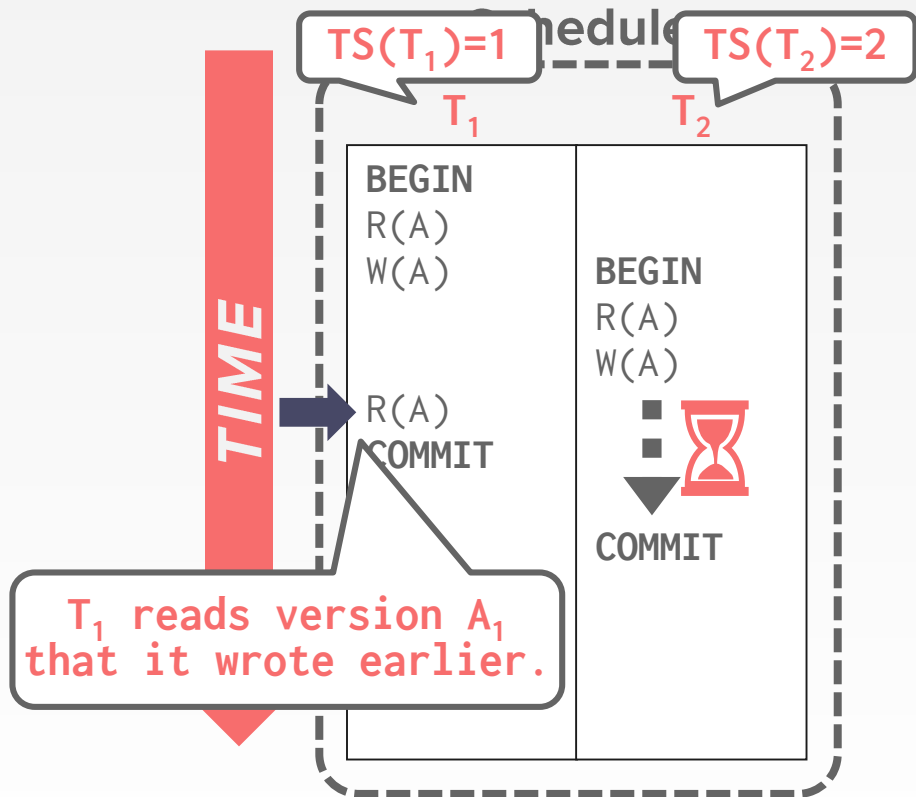
# MVCC – EXAMPLE #2



# MVCC – EXAMPLE #2



# MVCC – EXAMPLE #2

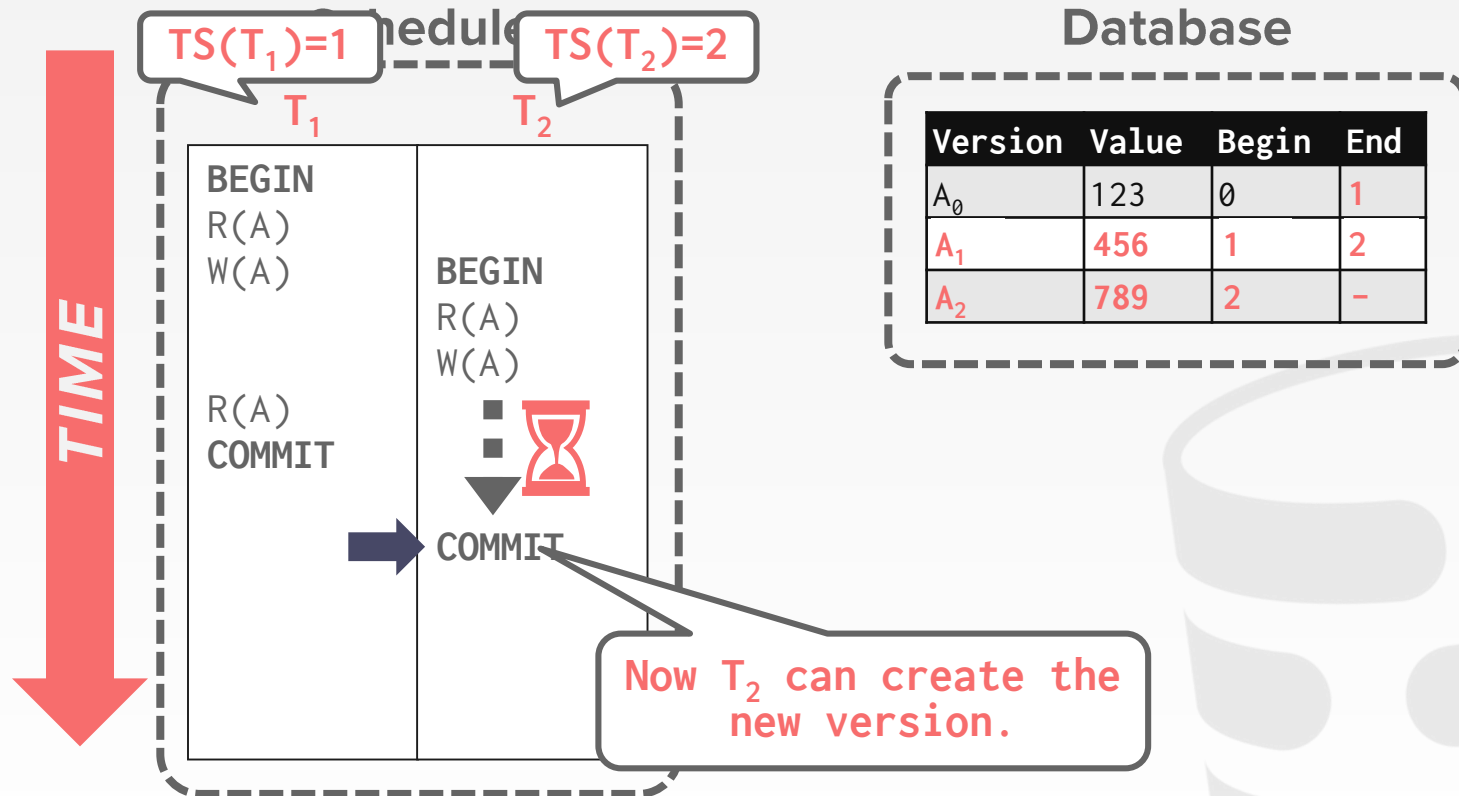


## Database

Version	Value	Begin	End
A <sub>0</sub>	123	0	1
A <sub>1</sub>	456	1	-



# MVCC – EXAMPLE #2



# MULTI-VERSION CONCURRENCY CONTROL

MVCC is more than just a concurrency control protocol. It completely affects how the DBMS manages transactions and the database.



# MVCC DESIGN DECISIONS

Concurrency Control Protocol

Version Storage

Garbage Collection

Index Management



# CONCURRENCY CONTROL PROTOCOL

## Approach #1: Timestamp Ordering

→ Assign txns timestamps that determine serial order.

## Approach #2: Optimistic Concurrency Control

→ Three-phase protocol from last class.

→ Use private workspace for new versions.

## Approach #3: Two-Phase Locking

→ Txns acquire appropriate lock on physical version before they can read/write a logical tuple.



# VERSION STORAGE

The DBMS uses the tuples' pointer field to create a **version chain** per logical tuple.

- This allows the DBMS to find the version that is visible to a particular txn at runtime.
- Indexes always point to the “head” of the chain.

Different storage schemes determine where/what to store for each version.



# VERSION STORAGE

## Approach #1: Append-Only Storage

→ New versions are appended to the same table space.

## Approach #2: Time-Travel Storage

→ Old versions are copied to separate table space.


## Approach #3: Delta Storage

→ The original values of the modified attributes are copied into a separate delta record space.



# APPEND-ONLY STORAGE

## Main Table

	KEY	VALUE	POINTER
$A_x$	XXX	\$111	
$A_{x+1}$	XXX	\$222	$\emptyset$
$B_x$	YYY	\$10	$\emptyset$

All of the physical versions of a logical tuple are stored in the same table space

On every update, append a new version of the tuple into an empty space in the table.

# APPEND-ONLY STORAGE

## Main Table

	KEY	VALUE	POINTER
$A_x$	XXX	\$111	●
$A_{x+1}$	XXX	\$222	∅
$B_x$	YYY	\$10	∅
$A_{x+2}$	XXX	\$333	∅

All of the physical versions of a logical tuple are stored in the same table space

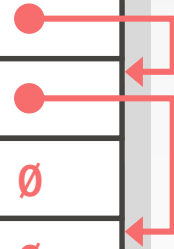
On every update, append a new version of the tuple into an empty space in the table.



# APPEND-ONLY STORAGE

## Main Table

	KEY	VALUE	POINTER
$A_x$	XXX	\$111	●
$A_{x+1}$	XXX	\$222	●
$B_x$	YYY	\$10	∅
$A_{x+2}$	XXX	\$333	∅



All of the physical versions of a logical tuple are stored in the same table space

On every update, append a new version of the tuple into an empty space in the table.

# VERSION CHAIN ORDERING

## Approach #1: Oldest-to-Newest (O2N)

- Just append new version to end of the chain.
- Have to traverse chain on look-ups.

## Approach #2: Newest-to-Oldest (N2O)

- Have to update index pointers for every new version.
- Don't have to traverse chain on look ups.



# TIME-TRAVEL STORAGE

*Main Table*

	KEY	VALUE	POINTER
$A_2$	XXX	\$222	●
$B_1$	YYY	\$10	

*Time-Travel Table*

	KEY	VALUE	POINTER
$A_1$	XXX	\$111	∅

On every update, copy the current version to the time-travel table. Update pointers.

# TIME-TRAVEL STORAGE

*Main Table*

	KEY	VALUE	POINTER
$A_2$	XXX	\$222	● →
$B_1$	YYY	\$10	

*Time-Travel Table*

	KEY	VALUE	POINTER
$A_1$	XXX	\$111	∅ ←
$A_2$	XXX	\$222	● ←

On every update, copy the current version to the time-travel table. Update pointers.

# TIME-TRAVEL STORAGE

## Main Table

	KEY	VALUE	POINTER
$A_3$	XXX	\$333	●
$B_1$	YYY	\$10	

## Time-Travel Table

	KEY	VALUE	POINTER
$A_1$	XXX	\$111	∅
$A_2$	XXX	\$222	●

On every update, copy the current version to the time-travel table. Update pointers.


Overwrite master version in the main table. Update pointers.

# DELTA STORAGE

## Main Table

	KEY	VALUE	POINTER
A <sub>1</sub>	XXX	\$111	
B <sub>1</sub>	YYY	\$10	

## Delta Storage Segment



On every update, copy only the values that were modified to the delta storage and overwrite the master version.

# DELTA STORAGE

## Main Table

	KEY	VALUE	POINTER
A <sub>1</sub>	XXX	\$111	
B <sub>1</sub>	YYY	\$10	

## Delta Storage Segment

	DELTA	POINTER
A <sub>1</sub>	(VALUE→\$111)	∅

On every update, copy only the values that were modified to the delta storage and overwrite the master version.

# DELTA STORAGE

*Main Table*

	KEY	VALUE	POINTER
A <sub>2</sub>	XXX	\$222	● →
B <sub>1</sub>	YYY	\$10	

*Delta Storage Segment*

	DELTA	POINTER
A <sub>1</sub>	(VALUE→\$111)	∅ ←
A <sub>2</sub>	(VALUE→\$222)	● ←

On every update, copy only the values that were modified to the delta storage and overwrite the master version.



# DELTA STORAGE

## Main Table

	KEY	VALUE	POINTER
A <sub>3</sub>	XXX	\$333	●
B <sub>1</sub>	YYY	\$10	

## Delta Storage Segment

	DELTA	POINTER
A <sub>1</sub>	(VALUE→\$111)	∅
A <sub>2</sub>	(VALUE→\$222)	●

On every update, copy only the values that were modified to the delta storage and overwrite the master version.

Txns can recreate old versions by applying the delta in reverse order.

# GARBAGE COLLECTION

The DBMS needs to remove reclaimable physical versions from the database over time.

- No active txn in the DBMS can “see” that version (SI).
- The version was created by an aborted txn.

Two additional design decisions:

- How to look for expired versions?
- How to decide when it is safe to reclaim memory?

# GARBAGE COLLECTION

## Approach #1: Tuple-level

- Find old versions by examining tuples directly.
- Background Vacuuming vs. Cooperative Cleaning

## Approach #2: Transaction-level

- Txns keep track of their old versions so the DBMS does not have to scan tuples to determine visibility.



# TUPLE-LEVEL GC

*Thread #1*

$T_{id}=12$

*Vacuum*



*Thread #2*

$T_{id}=25$

## Background Vacuuming:

Separate thread(s) periodically scan the table and look for reclaimable versions. Works with any storage.

	TXN-ID	BEGIN	END
$A_x$	0	1	9
$B_x$	0	1	9
$B_{x+1}$	0	10	20

# TUPLE-LEVEL GC

*Thread #1*

$T_{id}=12$

*Thread #2*

$T_{id}=25$

*Vacuum*



	TXN-ID	BEGIN	END
$A_x$	0	1	9
$B_x$	0	1	9
$B_{x+1}$	0	10	20

## Background Vacuuming:

Separate thread(s) periodically scan the table and look for reclaimable versions. Works with any storage.

# TUPLE-LEVEL GC

Thread #1

$T_{id}=12$

Thread #2

$T_{id}=25$

Vacuum



*Dirty?*

	TXN-ID	BEGIN	END
...			
$B_{x+1}$	0	10	20

## Background Vacuuming:

Separate thread(s) periodically scan the table and look for reclaimable versions. Works with any storage.


# TUPLE-LEVEL GC

*Thread #1*

$T_{id}=12$

*Thread #2*

$T_{id}=25$



## Background Vacuuming:

Separate thread(s) periodically scan the table and look for reclaimable versions. Works with any storage.

## Cooperative Cleaning:

Worker threads identify reclaimable versions as they traverse version chain. Only works with **O2N**.

# TUPLE-LEVEL GC

*Thread #1*

$T_{id}=12$

*Thread #2*

$T_{id}=25$



$A_{x+1}$

$A_{x+2}$

$A_{x+3}$

$B_x$

$B_{x+1}$

$B_{x+2}$

$B_{x+3}$

## Background Vacuuming:

Separate thread(s) periodically scan the table and look for reclaimable versions. Works with any storage.

## Cooperative Cleaning:

Worker threads identify reclaimable versions as they traverse version chain. Only works with **O2N**.



# TRANSACTION-LEVEL GC

Each txn keeps track of its read/write set.

The DBMS determines when all versions created by a finished txn are no longer visible.



# INDEX MANAGEMENT

PKey indexes always point to version chain head.


- How often the DBMS has to update the pkey index depends on whether the system creates new versions when a tuple is updated.
- If a txn updates a tuple's pkey attribute(s), then this is treated as an **DELETE** followed by an **INSERT**.

Secondary indexes are more complicated...

ARCHITECTURE

# WHY UBER ENGINEERING SWITCHED FROM POSTGRES TO MYSQL

JULY 26, 2016  
BY EVAN KLITZKE



# SECONDARY INDEXES

## Approach #1: Logical Pointers

- Use a fixed identifier per tuple that does not change.
- Requires an extra indirection layer.
- Primary Key vs. Tuple Id

## Approach #2: Physical Pointers

- Use the physical address to the version chain head.




# INDEX POINTERS



*PRIMARY INDEX*



*SECONDARY INDEX*




*Append-Only  
Newest-to-Oldest*

# INDEX POINTERS

GET(A) 

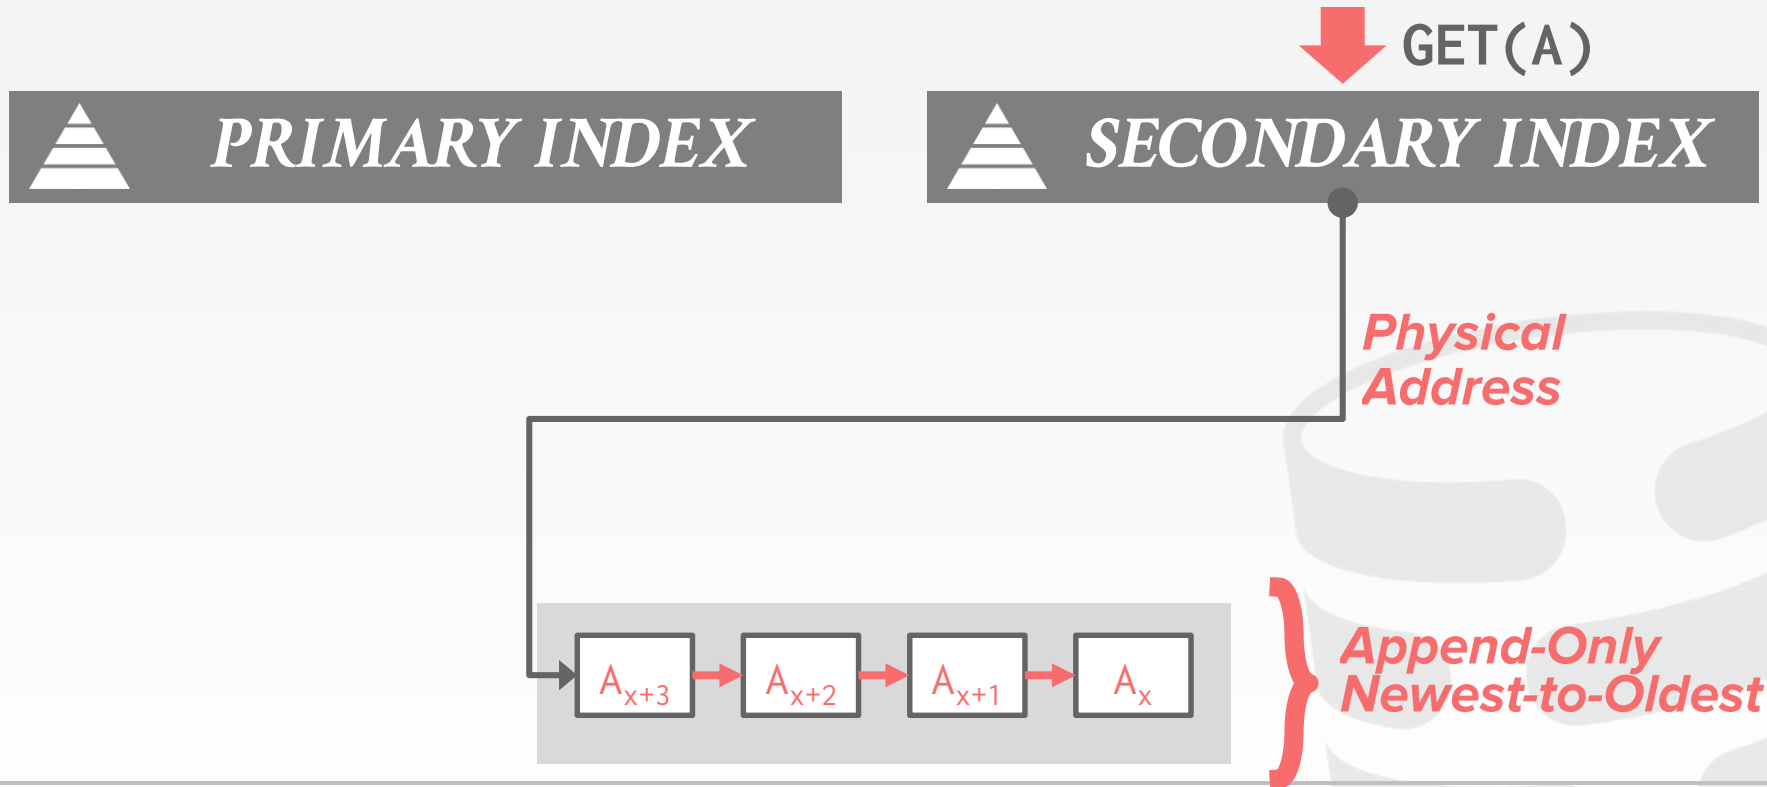


*Physical  
Address*

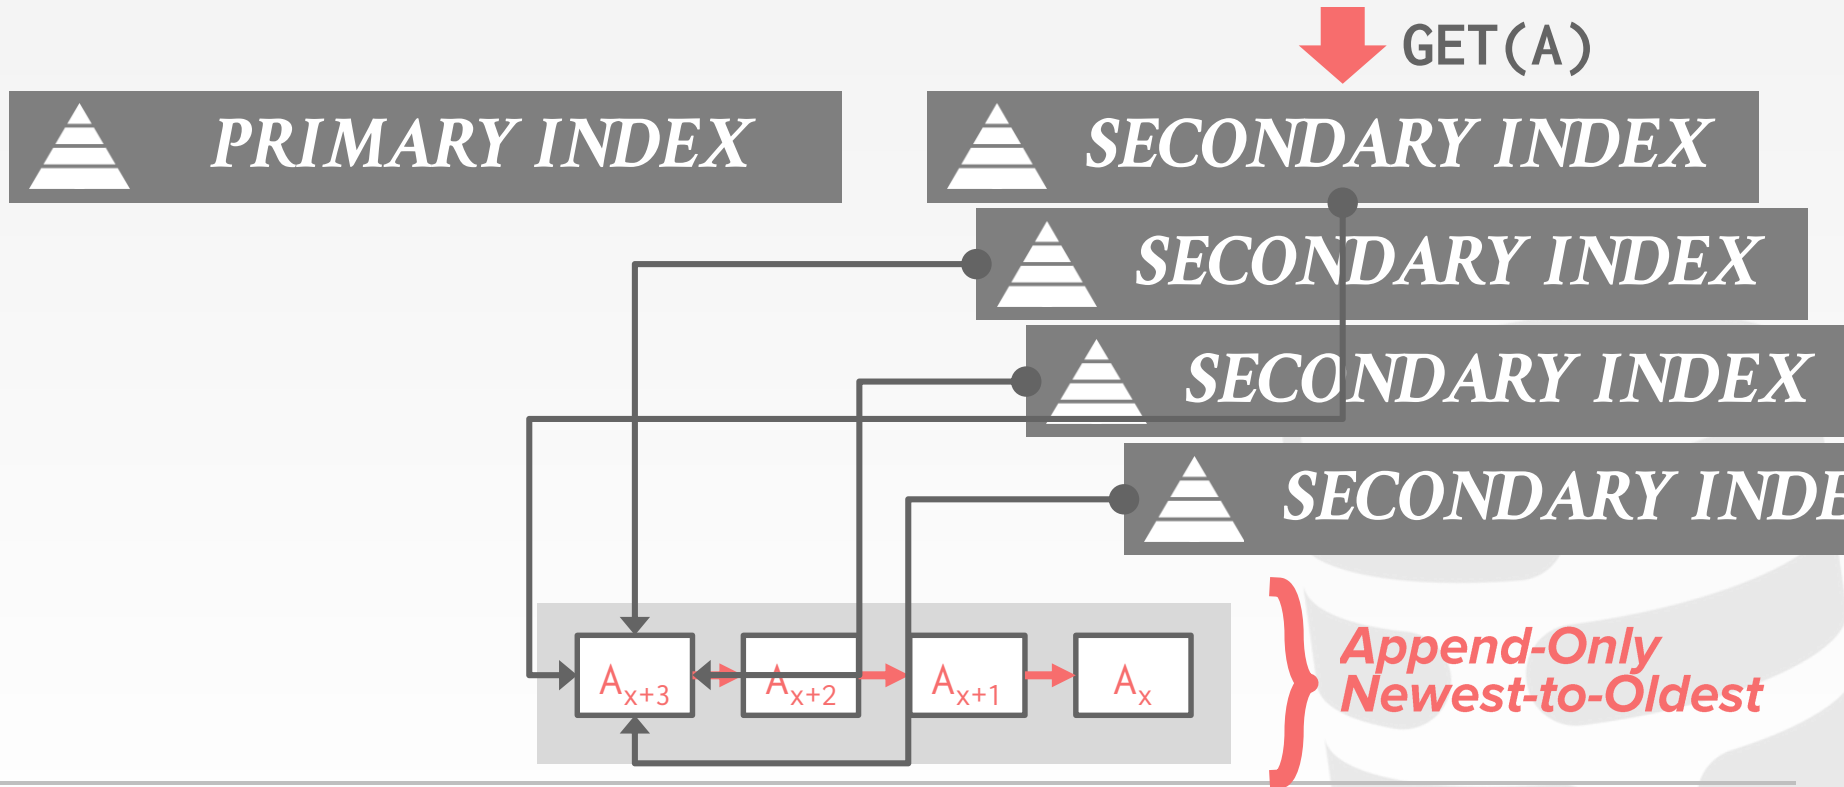


*Append-Only  
Newest-to-Oldest*

# INDEX POINTERS

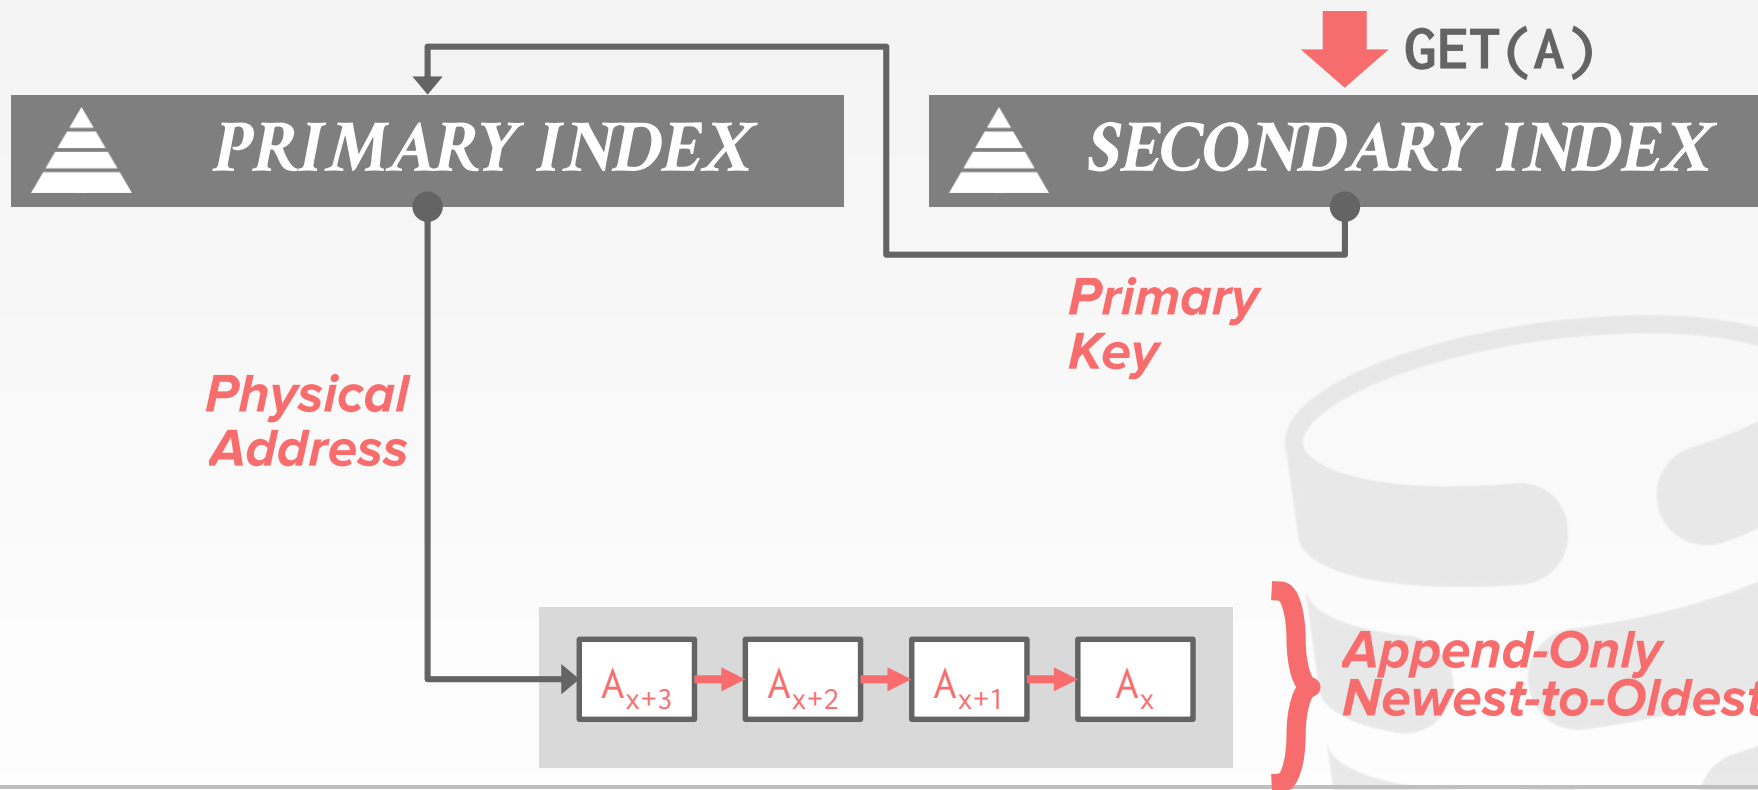


# INDEX POINTERS

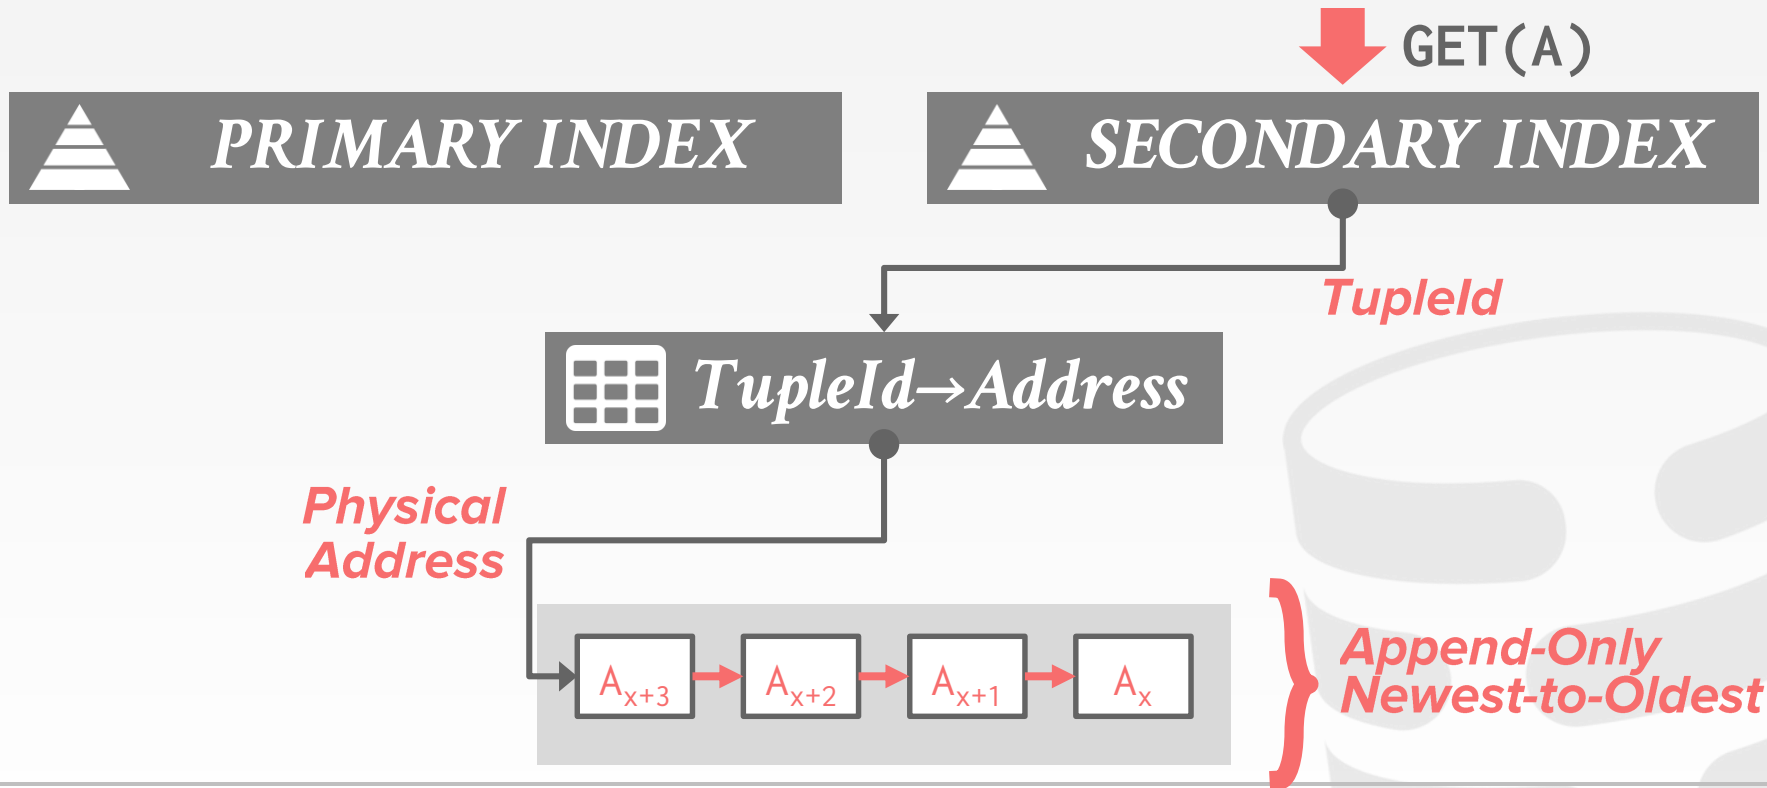




# INDEX POINTERS



# INDEX POINTERS



# MVCC IMPLEMENTATIONS

	<i>Protocol</i>	<i>Version Storage</i>	<i>Garbage Collection</i>	<i>Indexes</i>
Oracle	MV2PL	Delta	Vacuum	Logical
Postgres	MV-2PL/MV-TO	Append-Only	Vacuum	Physical
MySQL-InnoDB	MV-2PL	Delta	Vacuum	Logical
HYRISE	MV-OCC	Append-Only	-	Physical
Hekaton	MV-OCC	Append-Only	Cooperative	Physical
MemSQL	MV-OCC	Append-Only	Vacuum	Physical
SAP HANA	MV-2PL	Time-travel	Hybrid	Logical
NuoDB	MV-2PL	Append-Only	Vacuum	Logical
HyPer	MV-OCC	Delta	Txn-level	Logical

# CONCLUSION

MVCC is the widely used scheme in DBMSs. Even (NoSQL) systems that do not support multi-statement txns use it.



# NEXT CLASS

## Logging & Recovery

